

▼ Model Performances without Preprocessing

[Main Notebook Link](#)

▼ Installs

#Installs

```
!pip install pydot --quiet
!pip install gensim==3.8.3 --quiet
!pip install tensorflow-datasets --quiet
!pip install -U tensorflow-text==2.8.2 --quiet
!pip install transformers --quiet
!pip install datasets
!pip3 install emoji==0.6.0
```

```

23.4/23.4 MB 63.1 MB/s eta 0:00:00
Preparing metadata (setup.py) ... done
Building wheel for gensim (setup.py) ... done
4.9/4.9 MB 39.9 MB/s eta 0:00:00
498.1/498.1 MB 3.0 MB/s eta 0:00:00
42.6/42.6 kB 4.5 MB/s eta 0:00:00
5.8/5.8 MB 94.8 MB/s eta 0:00:00
1.1/1.1 MB 73.6 MB/s eta 0:00:00
462.3/462.3 kB 48.2 MB/s eta 0:00:00
1.4/1.4 MB 78.8 MB/s eta 0:00:00
4.9/4.9 MB 99.1 MB/s eta 0:00:00
7.0/7.0 MB 46.4 MB/s eta 0:00:00
7.8/7.8 MB 67.1 MB/s eta 0:00:00
200.1/200.1 kB 20.4 MB/s eta 0:00:00
Looking in indexes: https://pypi.org/simple, https://us-python.pkg.dev/colab-wheels/public/simple/
Collecting datasets
  Downloading datasets-2.11.0-py3-none-any.whl (468 kB)
    468.7/468.7 kB 9.0 MB/s eta 0:00:00
Requirement already satisfied: huggingface-hub<1.0.0,>=0.11.0 in /usr/local/lib/python3.9/dist-packages (from datasets) (0.11.0)
Requirement already satisfied: pyarrow>=8.0.0 in /usr/local/lib/python3.9/dist-packages (from datasets) (9.0.0)
Collecting multiprocessing
  Downloading multiprocessing-0.70.14-py39-none-any.whl (132 kB)
    132.9/132.9 kB 19.4 MB/s eta 0:00:00
Requirement already satisfied: packaging in /usr/local/lib/python3.9/dist-packages (from datasets) (23.0)
Collecting aiohttp
  Downloading aiohttp-3.8.4-cp39-cp39-manylinux_2_17_x86_64.manylinux2014_x86_64.whl (1.0 MB)
    1.0/1.0 MB 42.0 MB/s eta 0:00:00
Requirement already satisfied: numpy>=1.17 in /usr/local/lib/python3.9/dist-packages (from datasets) (1.22.4)
Requirement already satisfied: fsspec[http]>=2021.11.1 in /usr/local/lib/python3.9/dist-packages (from datasets) (2023.4.0)
Requirement already satisfied: pyyaml>=5.1 in /usr/local/lib/python3.9/dist-packages (from datasets) (6.0)
Collecting responses<0.19
  Downloading responses-0.18.0-py3-none-any.whl (38 kB)
    38.0/38.0 kB 15.5 MB/s eta 0:00:00
Requirement already satisfied: tqdm>=4.62.1 in /usr/local/lib/python3.9/dist-packages (from datasets) (4.65.0)
Requirement already satisfied: pandas in /usr/local/lib/python3.9/dist-packages (from datasets) (1.5.3)
Requirement already satisfied: requests>=2.19.0 in /usr/local/lib/python3.9/dist-packages (from datasets) (2.27.1)
Collecting xxhash
  Downloading xxhash-3.2.0-cp39-cp39-manylinux_2_17_x86_64.manylinux2014_x86_64.whl (212 kB)
    212.2/212.2 kB 26.6 MB/s eta 0:00:00
Collecting dill<0.3.7,>=0.3.0
  Downloading dill-0.3.6-py3-none-any.whl (110 kB)
    110.5/110.5 kB 15.5 MB/s eta 0:00:00
Collecting async-timeout<5.0,>=4.0.0a3
  Downloading async_timeout-4.0.2-py3-none-any.whl (5.8 kB)
Collecting frozenlist>=1.1.1
  Downloading frozenlist-1.3.3-cp39-cp39-manylinux_2_5_x86_64.manylinux1_x86_64.manylinux2_17_x86_64.manylinux2014_x86_64.whl (158.8 kB)
    158.8/158.8 kB 20.4 MB/s eta 0:00:00
Collecting yarl<2.0,>=1.0
  Downloading yarl-1.8.2-cp39-cp39-manylinux_2_17_x86_64.manylinux2014_x86_64.whl (264 kB)
    264.6/264.6 kB 28.1 MB/s eta 0:00:00
Requirement already satisfied: attrs>=17.3.0 in /usr/local/lib/python3.9/dist-packages (from aiohttp->datasets) (22.2.0)
Collecting multidict<7.0,>=4.5
  Downloading multidict-6.0.4-cp39-cp39-manylinux_2_17_x86_64.manylinux2014_x86_64.whl (114 kB)
    114.2/114.2 kB 4.5 MB/s eta 0:00:00
Requirement already satisfied: charset-normalizer<4.0,>=2.0 in /usr/local/lib/python3.9/dist-packages (from aiohttp->datasets) (3.2.0)
Collecting aiosignal>=1.1.2
  Downloading aiosignal-1.3.1-py3-none-any.whl (7.6 kB)
Requirement already satisfied: typing-extensions>=3.7.4.3 in /usr/local/lib/python3.9/dist-packages (from huggingface-hub<1.0.0,>=0.11.0->datasets) (4.5.0)
Requirement already satisfied: filelock in /usr/local/lib/python3.9/dist-packages (from huggingface-hub<1.0.0,>=0.11.0->datasets) (3.12.2)
```

▼ Imports

```
import numpy as np
import tensorflow as tf
from tensorflow import keras

from tensorflow.keras.layers import Embedding, Input, Dense, Lambda
from tensorflow.keras.models import Model
import tensorflow.keras.backend as K
import tensorflow_datasets as tfds
import tensorflow_text as tf_text

from transformers import BertTokenizer, TFBertModel, AutoModel, AutoTokenizer, TFAutoModelForSequenceClassification, TFAutoModel,

import nltk
from nltk.corpus import reuters
from nltk.data import find
from nltk.tokenize import RegexpTokenizer
from nltk.tokenize import wordpunct_tokenize
from nltk.tokenize import TweetTokenizer

import sklearn as sk
import os

import matplotlib.pyplot as plt

import re

#This continues to work with gensim 3.8.3. It doesn't yet work with 4.x.
#Make sure your pip install command specifies gensim==3.8.3
import gensim

from sklearn.metrics import classification_report
import seaborn as sns
```

▼ DataSet Import

```
import datasets
dataset = datasets.load_dataset('ucberkeley-dlab/measuring-hate-speech', 'binary')
df = dataset['train'].to_pandas()

##Citation
##@article{kennedy2020constructing,
  #title={Constructing interval variables via faceted Rasch measurement and multitask deep learning: a hate speech application},
  #author={Kennedy, Chris J and Bacon, Geoff and Sahn, Alexander and von Vacano, Claudia},
  #journal={arXiv preprint arXiv:2009.10277},
  #year={2020}
#}

Downloading readme: 100% 4.03k/4.03k [00:00<00:00, 250kB/s]
Downloading and preparing dataset parquet/ucberkeley-dlab--measuring-hate-speech to /root/.cache/huggingface/datasets/ucberke
Downloading data files: 100% 1/1 [00:01<00:00, 1.49s/it]

Downloading data: 100% 14.1M/14.1M [00:00<00:00, 28.1MB/s]

Extracting data files: 100% 1/1 [00:00<00:00, 25.07it/s]
Dataset parquet downloaded and prepared to /root/.cache/huggingface/datasets/ucberkeley-dlab___parquet/ucberkeley-dlab--measi
100% 1/1 [00:00<00:00, 15.64it/s]

# list of all the columns in the dataframe
for i in df.columns:
    print(i)

    comment_id
    annotator_id
    platform
    sentiment
    respect
    insult
```

```

humiliate
status
dehumanize
violence
genocide
attack_defend
hatespeech
hate_speech_score
text
infitms
outfitms
annotator_severity
std_err
annotator_infitms
annotator_outfitms
hypothesis
target_race_asian
target_race_black
target_race_latinx
target_race_middle_eastern
target_race_native_american
target_race_pacific_islander
target_race_white
target_race_other
target_race
target_religion_atheist
target_religion_buddhist
target_religion_christian
target_religion_hindu
target_religion_jewish
target_religion_mormon
target_religion_muslim
target_religion_other
target_religion
target_origin_immigrant
target_origin_migrant_worker
target_origin_specific_country
target_origin_undocumented
target_origin_other
target_origin
target_gender_men
target_gender_non_binary
target_gender_transgender_men
target_gender_transgender_unspecified
target_gender_transgender_women
target_gender_women
target_gender_other
target_gender
target_sexuality_bisexual
target_sexuality_gay
target_sexuality_lesbian
target_sexuality_straight

# creating arrays of text column and labels column for viewing and making text columns str
df['text'].astype(str)
df = df[['text', 'hate_speech_score']].groupby('text').mean().reset_index()
text = np.array(df['text'])
labels = np.array(df['hate_speech_score'])

```


▼ EDA

```

##EDA
import seaborn as sns
def str_labler(arr):
    new_arr = []
    for i in arr:
        if i >= .5:
            new_arr.append("Hateful")
        elif i < .5 and i > -1:
            new_arr.append("Neutral")
        else:
            new_arr.append("Supportive")
    return new_arr
labels_str = str_labler(labels)
sns.histplot(labels_str)

```

<Axes: ylabel='Count'>



```
supp_count = len([i for i in labels_str if i == "Supportive"])
hate_count = len([i for i in labels_str if i == "Hateful"])
neut_count = len([i for i in labels_str if i == "Neutral"])
total_count = len(labels_str)
print("Supportive Count: ", supp_count)
print("Supportive Prop: ", supp_count/total_count)
print("Hateful Count: ", hate_count)
print("Hateful Prop: ", hate_count/total_count)
print("Neutral Count: ", neut_count)
print("Neutral Prop: ", neut_count/total_count)
```

```
Supportive Count: 17553
Supportive Prop: 0.4436496903829142
Hateful Count: 10437
Hateful Prop: 0.2637937571085555
Neutral Count: 11575
Neutral Prop: 0.29255655250853024
```

```
supp_exs = df[df['hate_speech_score'] < -1].sort_values('hate_speech_score')
hate_exs = df[df['hate_speech_score'] > .5].sort_values('hate_speech_score', ascending = False)
neut_exs = df[df['hate_speech_score'].between(-1,.3)]
hate_neut_edge_exs = df[df['hate_speech_score'].between(.3,.7)]
```

supp_exs

hate_speech_score	text	2class_hate_speech
-------------------	------	--------------------

hate_exs

hate_speech_score	text	2class_hate_speech
0	! thank u! im transmasc and generally present ...	0
22476	My masculine is awakening from my spiritual el...	0
22484	My masculine is awakening through my spiritual...	0
22483	My masculine is awakening my from my spiritual...	0
22482	My masculine is awakening from my spiritual up...	0
...
19468	Islam needs to be fucking destroyed.	1
19467	Islam kills gays. Nothing to do with you being...	1
12330	Fuck yourself bitch	1
19462	Islam is bullsh#t.. a butch a beating, butcher...	1
18431	If anything, Fox news employs very lovely wome...	1

27990 rows x 3 columns

```
neut_exs[neut_exs['hate_speech_score'].between(-.01,.01)]
```

	hate_speech_score	text	2class_hate_speech
8	0	" Fuck you niggas ion give a fuck no more nigg...	0
10	0	" I got women tryna sneak me through there doo...	0
11	0	" I wish ur crack mom aborted you" well guess ...	0
12	0	" Pray for him" ha! The filthy traitor would b...	0
15	0	" is dat so hard for yu niggas to do😏 y'all wa...	0
...
39543	0	🔥 Vandals destroy Christian cemetery in Israel ...	0
39546	0	😭😭😭 so litt sus with the right bitch tho none o...	0
39553	0	😏 shhhhh them yung buls always wanna give me t...	0
39557	0	😏 this nigga shot out	0
39561	0	🎉 CONGRATULATIONS #MARYLAND #BALTIMORE LIBERALS...	0

11575 rows x 3 columns

```
hate_neut_edge_exs.sort_values('hate_speech_score', ascending = False)
```

hate_speech_score	text	2class_hate_speech
-------------------	------	--------------------

▼ Pre-Processing

```
def labler(arr):
    new_arr = []
    for i in arr:
        if i >= .5:
            new_arr.append(1)
        elif i < .5 and i > -1:
            new_arr.append(0)
        else:
            new_arr.append(2)
    return new_arr

#https://www.educative.io/answers/remove-all-the-punctuation-marks-from-a-sentence-using-regex

#labels = labler(labels)
new_df = df.apply(lambda x: labler(x) if x.name == 'hate_speech_score' else x,
                  axis = 0)

new_df['hate_speech_score'].value_counts()

2    17553
0     11575
1     10437
Name: hate_speech_score, dtype: int64

new_df['2class_hate_speech'] = [0 if i in [0,2] else 1 for i in new_df['hate_speech_score'].values]

new_df['2class_hate_speech'].value_counts()

0     29128
1     10437
Name: 2class_hate_speech, dtype: int64

np.random.seed(6543)
grouped = new_df[['hate_speech_score', 'text', '2class_hate_speech']].groupby('hate_speech_score', as_index = False).apply(lambda
df = grouped[['hate_speech_score', 'text', '2class_hate_speech']]
df
```

<ipython-input-14-aefeba8d0452>:2: FutureWarning: Not prepending group keys to the result index of transform-like apply. In 1 To preserve the previous behavior, use

```
>>> .groupby(..., group_keys=False)
```

To adopt the future behavior and silence this warning, use

```
>>> .groupby(..., group_keys=True)
grouped = new_df[['hate_speech_score', 'text', '2class_hate_speech']].groupby('hate_speech_score', as_index = False).apply(
```

hate_speech_score		text	2class_hate_speech	
0	2	! thank u! im transmasc and generally present ...	0	
1	1	!Go fuck yourself faggot!	1	
2	2	!flair [I love women and minorities]	0	
3	1	!flair [death to all niggers and gays]	1	
4	2	" 'convoluted' genealogy of Jesus"; was that c...	0	
...	
39560	1	🗨️ The limousine liberal Jews and WASPs who wa...	1	
39561	0	🗨️ CONGRATULATIONS #MARYLAND #BALTIMORE LIBERALS...	0	
39562	1	🗨️ another bitch will when another bitch can't	1	
39563	2	🗨️ next time stay your ass in the car before o...	0	
39564	2	🗨️ *IM SMARTER THAN SSSNIPERWOLF.I KNOW EVERY...	0	

► Train/Test/Val Split

```
[ ] ↳ 6 cells hidden
```

▼ Regular Tokenization

```
tokenizer = tf_text.WhitespaceTokenizer()
train_tokens = tokenizer.tokenize(train_examples)
val_tokens = tokenizer.tokenize(val_examples)
test_tokens = tokenizer.tokenize(test_examples)
```

```
max([len(i) for i in train_tokens.numpy()])
```

128

```
np.mean([len(i) for i in train_tokens.numpy()])
```

24.54015873015873

```
np.median([len(i) for i in train_tokens.numpy()])
```

19.0

▼ Baseline

```
def baseline_model(lbls):
    guesses = []
    for i in range(len(lbls)):
        x = np.random.uniform(low=0, high=1)
        if x <= supp_count/total_count:
            guesses.append(2)
        elif x <= (supp_count + hate_count)/total_count and x > supp_count/total_count:
            guesses.append(1)
        elif x > (supp_count + hate_count)/total_count and x<=1:
            guesses.append(0)
    return guesses
```

```
baseline_guesses = baseline_model(test_labels)
```

```
baseline_accuracy = np.mean([1 if baseline_guesses[i] == test_labels[i] else 0 for i in range(len(test_labels))])
baseline_accuracy

0.3486099410278012
```


```
print(classification_report(test_labels, baseline_guesses))
```

	precision	recall	f1-score	support
0	0.30	0.30	0.30	1750
1	0.26	0.26	0.26	1576
2	0.43	0.44	0.44	2609
accuracy			0.35	5935
macro avg	0.33	0.33	0.33	5935
weighted avg	0.35	0.35	0.35	5935

```
cm = tf.math.confusion_matrix(test_labels, baseline_guesses)
cm = cm/cm.numpy().sum(axis=1)[:, tf.newaxis]
```

```
plt.figure(figsize=(20,7))
sns.heatmap(
    cm, annot=True,
    xticklabels=["Neutral", "Hateful", "Supportive"],
    yticklabels=["Neutral", "Hateful", "Supportive"])
plt.xlabel("Predicted")
plt.ylabel("True")
```

```
Text(220.7222222222223, 0.5, 'True')
```



▼ Bert Tokenization

```
#Train and Test for Bert Based Models
bert_tokenizer = BertTokenizer.from_pretrained('bert-base-cased')
bert_model = TFBertModel.from_pretrained('bert-base-cased')

MAX_SEQUENCE_LENGTH = 128

bert_train_tokenized = bert_tokenizer(list(train_examples),
    max_length=MAX_SEQUENCE_LENGTH,
    truncation=True,
    padding='max_length',
    return_tensors='tf')
bert_train_inputs = [bert_train_tokenized.input_ids,
```

```

        bert_train_tokenized.token_type_ids,
        bert_train_tokenized.attention_mask]
bert_train_labels = np.array(train_labels)

bert_val_tokenized = bert_tokenizer(list(val_examples),
    max_length=MAX_SEQUENCE_LENGTH,
    truncation=True,
    padding='max_length',
    return_tensors='tf')
bert_val_inputs = [bert_val_tokenized.input_ids,
    bert_val_tokenized.token_type_ids,
    bert_val_tokenized.attention_mask]
bert_val_labels = np.array(val_labels)

bert_test_tokenized = bert_tokenizer(list(test_examples),
    max_length=MAX_SEQUENCE_LENGTH,
    truncation=True,
    padding='max_length',
    return_tensors='tf')
bert_test_inputs = [bert_test_tokenized.input_ids,
    bert_test_tokenized.token_type_ids,
    bert_test_tokenized.attention_mask]
bert_test_labels = np.array(test_labels)

```

```

Downloading (...)solve/main/vocab.txt: 100%          213k/213k [00:00<00:00, 3.94MB/s]

Downloading (...)okenizer_config.json: 100%          29.0/29.0 [00:00<00:00, 1.35kB/s]

Downloading (...)lve/main/config.json: 100%          570/570 [00:00<00:00, 25.4kB/s]

Downloading tf_model.h5: 100%                        527M/527M [00:07<00:00, 79.4MB/s]

```

Some layers from the model checkpoint at bert-base-cased were not used when initializing TFBertModel: ['nsp__cls', 'mlm__cls']
- This IS expected if you are initializing TFBertModel from the checkpoint of a model trained on another task or with another vocabulary.
- This IS NOT expected if you are initializing TFBertModel from the checkpoint of a model that you expect to be exactly identical to the model the checkpoint was trained on.
All the layers of TFBertModel were initialized from the model checkpoint at bert-base-cased.
If your task is similar to the task the model of the checkpoint was trained on, you can already use TFBertModel for predictions without further training.

```

#max(len(x) for x in bert_train_inputs)
np.unique(bert_train_labels[:5000], return_counts=True)

(array([0, 1, 2]), array([1434, 1340, 2226]))

```

▼ BERT CLS

```

#BERT Base Case
def create_bert_cls_classification_model(max_sequence_length=MAX_SEQUENCE_LENGTH,
    dropout = 0.3,
    hidden_size = 100,
    learning_rate=0.00001):
    """
    Build a classification model with BERT, where you apply CNN layers to the BERT output
    """

    bert_model.trainable = True
    max_length = MAX_SEQUENCE_LENGTH
    input_ids = tf.keras.layers.Input(shape=(max_length,), dtype=tf.int64, name='input_ids_layer')
    token_type_ids = tf.keras.layers.Input(shape=(max_length,), dtype=tf.int64, name='token_type_ids_layer')
    attention_mask = tf.keras.layers.Input(shape=(max_length,), dtype=tf.int64, name='attention_mask_layer')

    bert_inputs = {'input_ids': input_ids,
        'token_type_ids': token_type_ids,
        'attention_mask': attention_mask}

    bert_out = bert_model(bert_inputs)

    cls_token = bert_out[0][:, 0, :]

    hidden = tf.keras.layers.Dense(hidden_size, activation='relu', name='hidden_layer')(cls_token)

    hidden = tf.keras.layers.Dropout(dropout)(hidden)

    classification = tf.keras.layers.Dense(3, activation='softmax', name='classification_layer', kernel_regularizer='l1')(hidden)

```



```

classification_model = tf.keras.Model(inputs=[input_ids, token_type_ids, attention_mask], outputs=[classification])

classification_model.compile(optimizer=tf.keras.optimizers.Adam(learning_rate=learning_rate),
                             loss=tf.keras.losses.SparseCategoricalCrossentropy(from_logits=False),
                             metrics='accuracy')

return classification_model

bert_cls_classification_model = create_bert_cls_classification_model()
bert_cls_classification_model.summary()
#confirm all layers are frozen
bert_cls_classification_model_history = bert_cls_classification_model.fit(
    [bert_train_inputs[0], bert_train_inputs[1], bert_train_inputs[2]],
    bert_train_labels,
    validation_data=(bert_val_inputs[0], bert_val_inputs[1], bert_val_inputs[2]), bert_val_labels),
    batch_size=8,
    epochs=2)

Model: "model"

```

Layer (type)	Output Shape	Param #	Connected to
attention_mask_layer (InputLayer)	[(None, 128)]	0	[]
input_ids_layer (InputLayer)	[(None, 128)]	0	[]
token_type_ids_layer (InputLayer)	[(None, 128)]	0	[]
tf_bert_model (TFBertModel)	TFBaseModelOutputWithPoolingAndCrossAttentions(last_hidden_state=(None, 128, 768), pooler_output=(None, 768), past_key_values=None, hidden_states=None, attentions=None, cross_attentions=None)	108310272	['attention_mask_layer[0][0]', 'input_ids_layer[0][0]', 'token_type_ids_layer[0][0]']
tf.__operators__.getitem (SlicingOpLambda)	(None, 768)	0	['tf_bert_model[0][0]']
hidden_layer (Dense)	(None, 100)	76900	['tf.__operators__.getitem[0][0]']
dropout_37 (Dropout)	(None, 100)	0	['hidden_layer[0][0]']
classification_layer (Dense)	(None, 3)	303	['dropout_37[0][0]']

```

=====
Total params: 108,387,475
Trainable params: 108,387,475
Non-trainable params: 0
Epoch 1/2
WARNING:tensorflow:Gradients do not exist for variables ['tf_bert_model/bert/pooler/dense/kernel:0', 'tf_bert_model/bert/pooler/dense/kernel:0', 'tf_bert_model/bert/pooler/dense/kernel:0']
WARNING:tensorflow:Gradients do not exist for variables ['tf_bert_model/bert/pooler/dense/kernel:0', 'tf_bert_model/bert/pooler/dense/kernel:0', 'tf_bert_model/bert/pooler/dense/kernel:0']
3462/3462 [=====] - 825s 235ms/step - loss: 1.1745 - accuracy: 0.6077 - val_loss: 1.0605 - val_accuracy: 0.6077
Epoch 2/2
3462/3462 [=====] - 813s 235ms/step - loss: 0.9898 - accuracy: 0.6909 - val_loss: 1.0136 - val_accuracy: 0.6909

cls_predictions = bert_cls_classification_model.predict([bert_test_inputs[0], bert_test_inputs[1], bert_test_inputs[2]])
cls_predictions = tf.argmax(cls_predictions, axis=-1)

print(classification_report(bert_test_labels, cls_predictions))#, target_names=["Neutral", "Hateful", "Supportive"])

```


	precision	recall	f1-score	support
0	0.65	0.55	0.59	1750
1	0.73	0.84	0.78	1576
2	0.85	0.85	0.85	2609
accuracy			0.76	5935
macro avg	0.74	0.75	0.74	5935
weighted avg	0.76	0.76	0.76	5935

```

cm = tf.math.confusion_matrix(bert_test_labels, cls_predictions)
cm = cm/cm.numpy().sum(axis=1)[:, tf.newaxis]
plt.figure(figsize=(20,7))
sns.heatmap(
    cm, annot=True)
plt.xlabel("Predicted")
plt.ylabel("True")

```

Text(220.7222222222223, 0.5, 'True')



► Bert CNN

[] ↳ 8 cells hidden

▼ Combined Roberta Model

```

from transformers import RobertaConfig, TFRobertaModel

#Roberta Model, Tokenizer, and Train and Test Sets

roberta_tokenizer = AutoTokenizer.from_pretrained("roberta-base")
roberta_model = TFRobertaModel.from_pretrained("roberta-base")

MAX_SEQUENCE_LENGTH = 128
roberta_train_tokenized = roberta_tokenizer(list(train_examples),
    max_length=MAX_SEQUENCE_LENGTH,
    truncation=True,
    padding='max_length',
    add_special_tokens=True,
    return_tensors='tf')
print(roberta_train_tokenized)
roberta_train_inputs = [roberta_train_tokenized.input_ids,
    roberta_train_tokenized.attention_mask]
roberta_train_labels = np.array(train_labels)

roberta_val_tokenized = roberta_tokenizer(list(val_examples),
    max_length=MAX_SEQUENCE_LENGTH,
    truncation=True,
    padding='max_length',
    return_tensors='tf')
roberta_val_inputs = [roberta_val_tokenized.input_ids,

```

```

        roberta_val_tokenized.attention_mask]
roberta_val_labels = np.array(val_labels)

roberta_test_tokenized = roberta_tokenizer(list(test_examples),
        max_length=MAX_SEQUENCE_LENGTH,
        truncation=True,
        padding='max_length',
        return_tensors='tf')
roberta_test_inputs = [roberta_test_tokenized.input_ids,
        #roberta_test_tokenized.token_type_ids,
        roberta_test_tokenized.attention_mask]
roberta_test_labels = np.array(test_labels)

Downloading (...)lve/main/config.json: 100%                481/481 [00:00<00:00, 9.33kB/s]

Downloading (...)olve/main/vocab.json: 100%                899k/899k [00:00<00:00, 6.22MB/s]

Downloading (...)olve/main/merges.txt: 100%                456k/456k [00:00<00:00, 3.60MB/s]

Downloading (...)main/tokenizer.json: 100%                 1.36M/1.36M [00:00<00:00, 8.04MB/s]

Downloading tf_model.h5: 100%                             657M/657M [00:08<00:00, 96.1MB/s]

Some layers from the model checkpoint at roberta-base were not used when initializing TFRobertaModel: ['lm_head']
- This IS expected if you are initializing TFRobertaModel from the checkpoint of a model trained on another task or with anot
- This IS NOT expected if you are initializing TFRobertaModel from the checkpoint of a model that you expect to be exactly ic
All the layers of TFRobertaModel were initialized from the model checkpoint at roberta-base.
If your task is similar to the task the model of the checkpoint was trained on, you can already use TFRobertaModel for predic
{'input_ids': <tf.Tensor: shape=(27696, 128), dtype=int32, numpy=
array([[ 0, 1039, 6785, ..., 1, 1, 1],
       [ 0, 1039, 298, ..., 1, 1, 1],
       [ 0, 20328, 47, ..., 1, 1, 1],
       ...,
       [ 0, 597, 24029, ..., 1, 1, 1],
       [ 0, 15698, 286, ..., 1, 1, 1],
       [ 0, 37294, 219, ..., 1, 1, 1]], dtype=int32)>, 'attention_mask': <tf.Tensor: shape=(27696, 128), dt
array([[1, 1, 1, ..., 0, 0, 0],
       [1, 1, 1, ..., 0, 0, 0],
       [1, 1, 1, ..., 0, 0, 0],
       ...,
       [1, 1, 1, ..., 0, 0, 0],
       [1, 1, 1, ..., 0, 0, 0],
       [1, 1, 1, ..., 0, 0, 0]], dtype=int32)>}]

def create_robertal_classification_model(max_sequence_length=MAX_SEQUENCE_LENGTH,
        num_filters = [100, 100, 50, 25],
        kernel_sizes = [2, 3, 4, 5],
        dropout = 0.3,
        hidden_size = 100,
        learning_rate=.00001):
    """
    Build a classification model with BERT, where you apply CNN layers to the BERT output
    """

    roberta_model.trainable = True
    #max_length = 100
    input_ids = tf.keras.layers.Input(shape=(max_sequence_length,), dtype=tf.int64, name='input_ids_layer')
    # token_type_ids = tf.keras.layers.Input(shape=(max_sequence_length,), dtype=tf.int64, name='token_type_ids_layer')
    attention_mask = tf.keras.layers.Input(shape=(max_sequence_length,), dtype=tf.int64, name='attention_mask_layer')

    roberta_inputs = {'input_ids': input_ids,
        'attention_mask': attention_mask}

    roberta_out = roberta_model(roberta_inputs)

    pooler_token = roberta_out[1]
    print(pooler_token)
    hidden = tf.keras.layers.Dense(hidden_size, activation='relu', name='hidden_layer')(pooler_token)

    hidden = tf.keras.layers.Dropout(dropout)(hidden)

    classification = tf.keras.layers.Dense(3, activation='softmax', name='classification_layer')(hidden)

    classification_model = tf.keras.Model(inputs=[input_ids, attention_mask], outputs=[classification])

    classification_model.compile(optimizer=tf.keras.optimizers.Adam(learning_rate=learning_rate),
        loss=tf.keras.losses.SparseCategoricalCrossentropy(from_logits=False),
        metrics='accuracy')

```

```

return classification_model

robertal_classification_model = create_robertal_classification_model()
robertal_classification_model.summary()
#confirm all layers are frozen
robertal_classification_model_history = robertal_classification_model.fit(
    [roberta_train_inputs[0], roberta_train_inputs[1]],
    roberta_train_labels,
    validation_data=([roberta_test_inputs[0], roberta_test_inputs[1]], roberta_test_labels),
    batch_size=8,
    epochs=2)

KerasTensor(type_spec=TensorSpec(shape=(None, 768), dtype=tf.float32, name=None), name='tf_roberta_model/roberta/pooler/dense
Model: "model"

```

Layer (type)	Output Shape	Param #	Connected to
attention_mask_layer (InputLayer)	[(None, 128)]	0	[]
input_ids_layer (InputLayer)	[(None, 128)]	0	[]
tf_roberta_model (TFRobertaModel)	TFBaseModelOutputWithPoolingAndCrossAttentions(last_hidden_state=(None, 128, 768), pooler_output=(None, 768), past_key_values=None, hidden_states=None, attentions=None, cross_attentions=None)	124645632	['attention_mask_layer[0][0]', 'input_ids_layer[0][0]']
hidden_layer (Dense)	(None, 100)	76900	['tf_roberta_model[0][1]']
dropout_37 (Dropout)	(None, 100)	0	['hidden_layer[0][0]']
classification_layer (Dense)	(None, 3)	303	['dropout_37[0][0]']
Total params: 124,722,835			
Trainable params: 124,722,835			
Non-trainable params: 0			

```

Epoch 1/2
3462/3462 [=====] - 861s 246ms/step - loss: 0.7949 - accuracy: 0.6232 - val_loss: 0.6932 - val_accuracy: 0.5965
Epoch 2/2
3462/3462 [=====] - 851s 246ms/step - loss: 0.6859 - accuracy: 0.6833 - val_loss: 0.5965 - val_accuracy: 0.5965

```

```

predictions1 = robertal_classification_model.predict([roberta_test_inputs[0], roberta_test_inputs[1]])

predictions1 = tf.argmax(predictions1, axis=-1)

from sklearn.metrics import classification_report
print(classification_report(roberta_test_labels, predictions1), target_names=["Neutral", "Hateful", "Supportive"])

```

	precision	recall	f1-score	support
0	0.58	0.59	0.59	1750
1	0.70	0.83	0.76	1576
2	0.88	0.78	0.83	2609
accuracy			0.74	5935
macro avg	0.72	0.73	0.73	5935
weighted avg	0.75	0.74	0.74	5935

```

cm = tf.math.confusion_matrix(roberta_test_labels, predictions1)
cm = cm/cm.numpy().sum(axis=1)[:, tf.newaxis]

```

```


plt.figure(figsize=(20,7))
sns.heatmap(
    cm, annot=True,
    xticklabels=["Neutral", "Hateful", "Supportive"],

```

```

yticklabels=["Neutral", "Hateful", "Supportive"])
plt.xlabel("Predicted")
plt.ylabel("True")
Text(220.722222222223, 0.5, 'True')


```



```

from matplotlib import pyplot as plt
plt.plot(robertal_classification_model_history.history['accuracy'])
plt.plot(robertal_classification_model_history.history['val_accuracy'])
plt.title('model accuracy')
plt.ylabel('accuracy')
plt.xlabel('epoch')
plt.legend(['train', 'val'], loc='upper left')
plt.show()

```



```

MAX_SEQUENCE_LENGTH = 128
binary_roberta_train_tokenized = roberta_tokenizer(list(binary_train_examples),
max_length=MAX_SEQUENCE_LENGTH,
truncation=True,
padding='max_length',
return_tensors='tf')
binary_roberta_train_inputs = [binary_roberta_train_tokenized.input_ids,
binary_roberta_train_tokenized.attention_mask]

```

```

binary_roberta_train_labels = np.array(binary_train_labels)

binary_roberta_val_tokenized = roberta_tokenizer(list(binary_val_examples),
max_length=MAX_SEQUENCE_LENGTH,
truncation=True,
padding='max_length',
return_tensors='tf')
binary_roberta_val_inputs = [binary_roberta_val_tokenized.input_ids,
binary_roberta_val_tokenized.attention_mask]
binary_roberta_val_labels = np.array(binary_val_labels)

binary_roberta_test_tokenized = roberta_tokenizer(list(binary_test_examples),
max_length=MAX_SEQUENCE_LENGTH,
truncation=True,
padding='max_length',
return_tensors='tf')
binary_roberta_test_inputs = [binary_roberta_test_tokenized.input_ids,
binary_roberta_test_tokenized.attention_mask]
binary_roberta_test_labels = np.array(binary_test_labels)

def create_roberta2_classification_model(max_sequence_length=MAX_SEQUENCE_LENGTH,
num_classes = 1,
dropout = 0.3,
hidden_size = 100,
learning_rate=0.00001):
"""
Build a classification model with BERT, where you apply CNN layers to the BERT output
"""

roberta_model.trainable = True
#max_length = 100
input_ids = tf.keras.layers.Input(shape=(max_sequence_length,), dtype=tf.int64, name='input_ids_layer')
# token_type_ids = tf.keras.layers.Input(shape=(max_sequence_length,), dtype=tf.int64, name='token_type_ids_layer')
attention_mask = tf.keras.layers.Input(shape=(max_sequence_length,), dtype=tf.int64, name='attention_mask_layer')

roberta_inputs = {'input_ids': input_ids,
'attention_mask': attention_mask}

roberta_out = roberta_model(roberta_inputs)

cls_token = roberta_out[1]

hidden = tf.keras.layers.Dense(hidden_size, activation='relu', name='hidden_layer')(cls_token)

hidden = tf.keras.layers.Dropout(dropout)(hidden)

classification = tf.keras.layers.Dense(num_classes, activation='sigmoid', name='classification_layer')(hidden)

classification_model = tf.keras.Model(inputs=[input_ids, attention_mask], outputs=[classification])

classification_model.compile(optimizer=tf.keras.optimizers.Adam(learning_rate=learning_rate),
loss=tf.keras.losses.BinaryCrossentropy(from_logits=False),
metrics='accuracy')

return classification_model

binary_roberta2_classification_model = create_roberta2_classification_model()
binary_roberta2_classification_model.summary()
#confirm all layers are frozen
binary_roberta2_classification_model_history = binary_roberta2_classification_model.fit(
[binary_roberta_train_inputs[0], binary_roberta_train_inputs[1]],
binary_roberta_train_labels,
validation_data=([binary_roberta_val_inputs[0], binary_roberta_val_inputs[1]], binary_roberta_val_labels),
batch_size=8,
epochs=2)

Model: "model_2"

```

Layer (type)	Output Shape	Param #	Connected to
attention_mask_layer (InputLayer)	[(None, 128)]	0	[]
input_ids_layer (InputLayer)	[(None, 128)]	0	[]
tf_roberta_model (TFRobertaMod	TFBaseModelOutputWi	124645632	['attention_mask_layer[0][0]',

```

el)                                thPoolingAndCrossAt          'input_ids_layer[0][0]']
tentions(last_hidde
n_state=(None, 128,
768),
pooler_output=(Non
e, 768),
past_key_values=No
ne, hidden_states=N
one, attentions=Non
e, cross_attentions
=None)

hidden_layer (Dense)              (None, 100)              76900              ['tf_roberta_model[2][1]']
dropout_39 (Dropout)              (None, 100)              0                  ['hidden_layer[0][0]']
classification_layer (Dense)      (None, 1)                101                ['dropout_39[0][0]']

=====
Total params: 124,722,633
Trainable params: 124,722,633
Non-trainable params: 0

Epoch 1/2
3462/3462 [=====] - 878s 246ms/step - loss: 0.2316 - accuracy: 0.9032 - val_loss: 0.4643 - val_accu
Epoch 2/2
3462/3462 [=====] - 849s 245ms/step - loss: 0.1745 - accuracy: 0.9280 - val_loss: 0.5233 - val_accu

binary_predictions2 = binary_roberta2_classification_model.predict([binary_roberta_test_inputs[0], binary_roberta_test_inputs[1]])

binary_predictions2 = [1 if i[0] > .5 else 0 for i in binary_predictions2]

print(classification_report(binary_roberta_test_labels, binary_predictions2))#, target_names=["Neutral", "Hateful", "Supportive"])

      precision    recall  f1-score   support

     0       0.95      0.97      0.96      4359
     1       0.91      0.85      0.88      1576


 accuracy                   0.94      5935
 macro avg       0.93      0.91      0.92      5935
 weighted avg    0.94      0.94      0.94      5935


cm = tf.math.confusion_matrix(binary_roberta_test_labels, binary_predictions2)
cm = cm/cm.numpy().sum(axis=1)[:, tf.newaxis]

plt.figure(figsize=(20,7))
sns.heatmap(
    cm, annot=True,
    xticklabels=["Neutral", "Hateful", "Supportive"],
    yticklabels=["Neutral", "Hateful", "Supportive"])
plt.xlabel("Predicted")
plt.ylabel("True")

```

Text(220.7222222222223, 0.5, 'True')



```
combined_preds = []
for i in range(len(predictions1)):
    if predictions1[i] == 0 and binary_predictions2[i] == 1:
        combined_preds.append(binary_predictions2[i])
    else:
        combined_preds.append(predictions1[i])
combined_preds = np.array(combined_preds)
```



```
print(len(predictions1))
print(len(binary_predictions2))
```

5935
5935

True Predicted

```
print(classification_report(roberta_test_labels, combined_preds))
```

	precision	recall	f1-score	support
0	0.63	0.58	0.60	1750
1	0.72	0.92	0.81	1576
2	0.88	0.78	0.83	2609
accuracy			0.76	5935
macro avg	0.74	0.76	0.75	5935
weighted avg	0.76	0.76	0.76	5935


```
plt.figure(figsize=(20,7))
cm = tf.math.confusion_matrix(roberta_test_labels, combined_preds)
cm = cm/cm.numpy().sum(axis=1)[:, tf.newaxis]
```

<Figure size 2000x700 with 0 Axes>

```
plt.figure(figsize=(20,7))
sns.heatmap(
    cm, annot=True,
    xticklabels=["Neutral", "Hateful", "Supportive"],
    yticklabels=["Neutral", "Hateful", "Supportive"])
plt.xlabel("Predicted")
plt.ylabel("True")
```



```
Text(220.7222222222223, 0.5, 'True')
```



```
incorrects = []
for i in range(len(combined_preds)):
    if combined_preds[i] != roberta_test_labels[i]:
        prediction=combined_preds[i]
        label=roberta_test_labels[i]
        example = list(test_examples)[i]
        incorrects.append([prediction,label, example])
```

0

```
for i in range(10):
    print(f'Text {i}: ', incorrects[i][2])
    print(f'Predicted Label {i}: ', incorrects[i][0])
    print(f'True Label {i}: ', incorrects[i][1])
```

```
Text 0: It was used during beatings, rapes, and murders of black people
Predicted Label 0: 0
True Label 0: 2
Text 1: It's not a shit hole because that's where the white people live. That kid wasn't white.
Predicted Label 1: 2
True Label 1: 0
Text 2: Let me fuck you as your sucking your own pulsating cock
Predicted Label 2: 2
True Label 2: 0
Text 3: @TylerTrollbane @penelope7usa Liberals: White genocide NOW! "Conservatives": Hold on! We need a phase-in so we can /
Predicted Label 3: 0
True Label 3: 1
Text 4: After recent performance of India in icc tournament ( S.a smiles from a distance ) C H O K E R S
Predicted Label 4: 2
True Label 4: 0
Text 5: The concept of "gendered brains" is demonstrably false through the very science that those mainstream atheists claim
Predicted Label 5: 0
True Label 5: 2
Text 6: i really hate girls that are stuck up for no reason like damn sis.
Predicted Label 6: 0
True Label 6: 2
Text 7: I Love to wear my tiny #Cage and #locked my Small #dick #me #bottom #lockedboy #gayfaggot #faggot #TeamLocked URL
Predicted Label 7: 1
True Label 7: 0
Text 8: I need my fat light ass fucked by a big black nigger dick 🍆🔥 URL
Predicted Label 8: 0
True Label 8: 2
Text 9: Evangelicals are angry with Trump for using God's name to swear. God is angry that Evangelicals for using God's name
Predicted Label 9: 0
True Label 9: 2
```